Homework-8: TinyML using Edge Impulse

Q1.1

off dataset

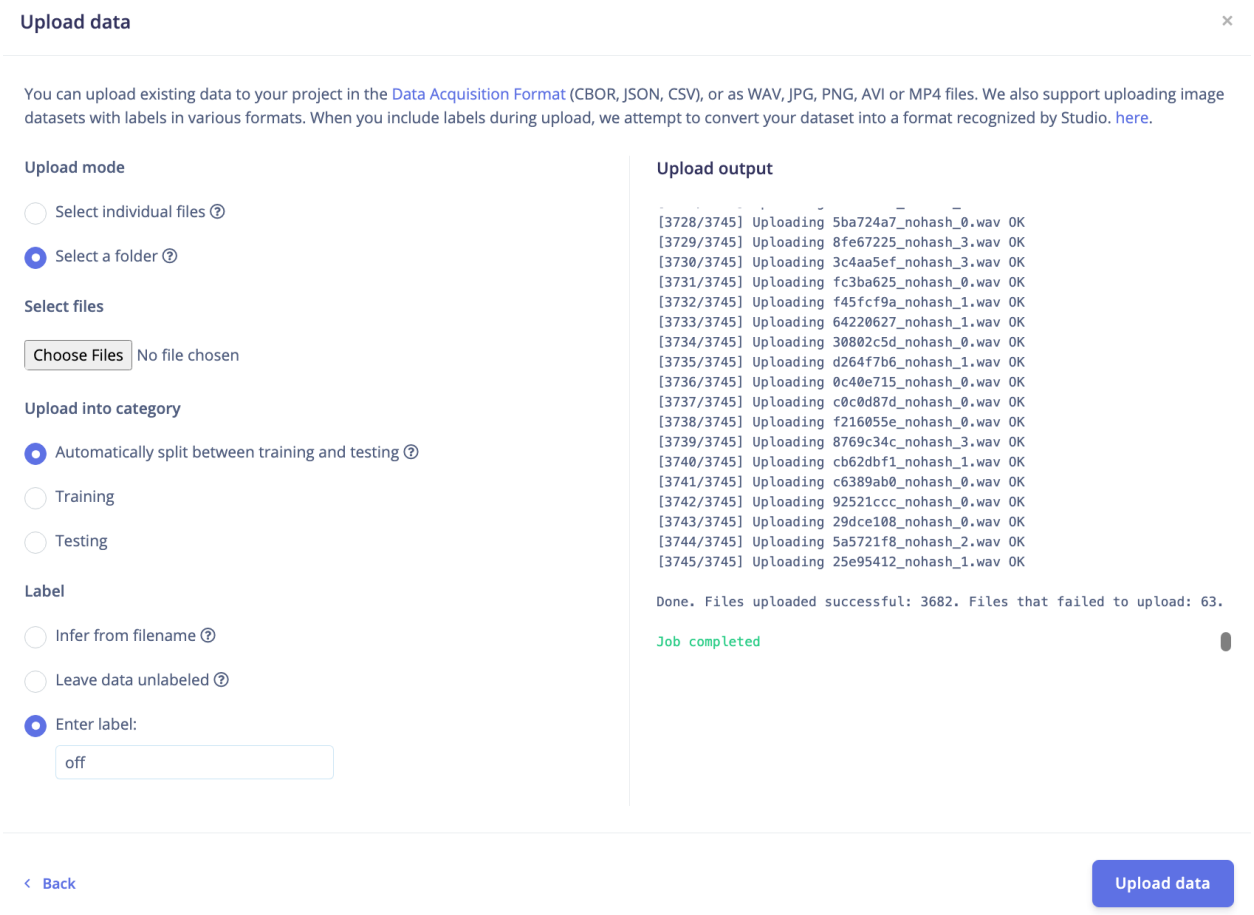

on dataset

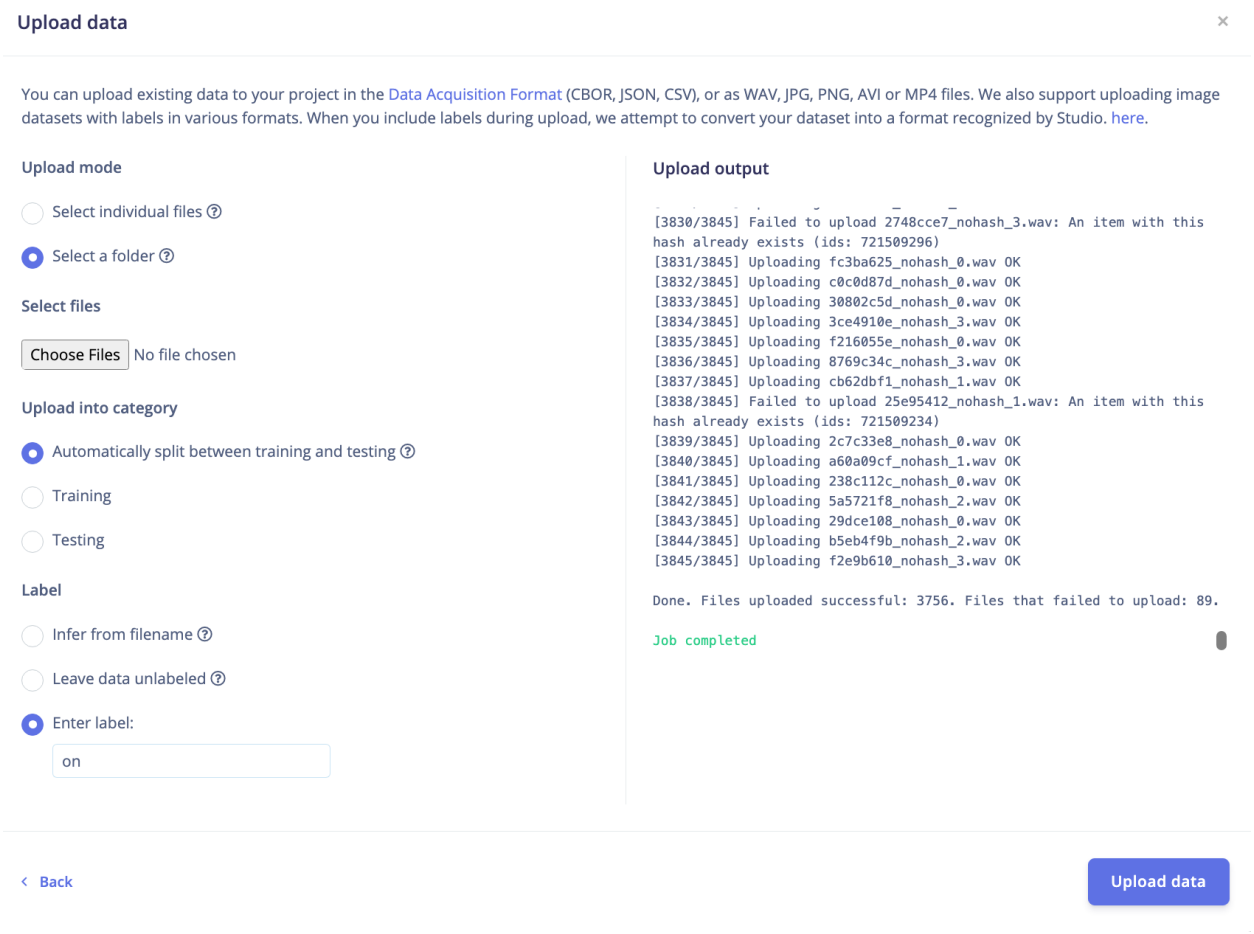
Upload data
You can upload existing data to your project in the Data Acquisition Format (CBOR, JSON, CSV), or as WAV, JPG, PNG, AVI or MP4 files. We also support uploading image datasets with labels in various formats. When you include labels during upload, we attempt to convert your dataset into a format recognized by Studio. here.
Upload mode
Select individual files
Select a folder
Select files
Choose Files
No file chosen
Upload into category
Automatically split between training and testing
Training
Testing
Label
Infer from filename
Leave data unlabeled
Enter label:
on
Upload output
[3830/3845] Failed to upload 2748cce7_nohash_3.wav: An item with this hash already exists (ids: 721509296)
[3831/3845] Uploading fc3ba625_nohash_0.wav OK
[3832/3845] Uploading c0c0d87d_nohash_0.wav OK
[3833/3845] Uploading 30802c5d_nohash_0.wav OK
[3834/3845] Uploading 3ce4910e_nohash_3.wav OK
[3835/3845] Uploading f216055e_nohash_0.wav OK
[3836/3845] Uploading 8769c34c_nohash_3.wav OK
[3837/3845] Uploading cb62dbf1_nohash_1.wav OK
[3838/3845] Failed to upload 25e95412_nohash_1.wav: An item with this hash already exists (ids: 721509234)
[3839/3845] Uploading 2c7c33e8_nohash_0.wav OK
[3840/3845] Uploading a60a09cf_nohash_1.wav OK
[3841/3845] Uploading 238c112c_nohash_0.wav OK
[3842/3845] Uploading 5a5721f8_nohash_2.wav OK
[3843/3845] Uploading 29dce108_nohash_0.wav OK
[3844/3845] Uploading b5eb4f9b_nohash_2.wav OK
[3845/3845] Uploading f2e9b610_nohash_3.wav OK
Done. Files uploaded successful: 3756. Files that failed to upload: 89.
Job completed
Back
Upload data

Others dataset

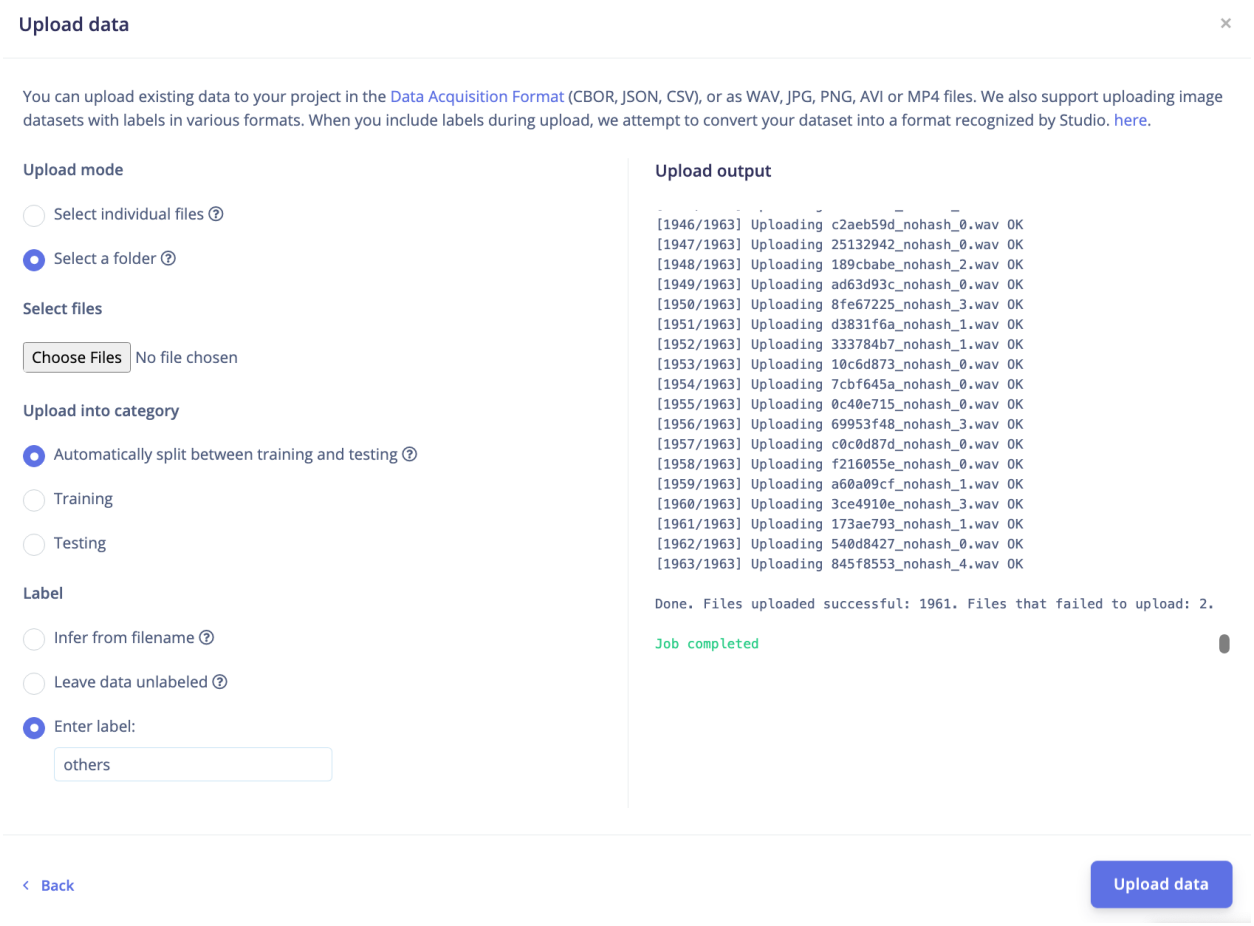

Silent dataset

Upload data

You can upload existing data to your project in the Data Acquisition Format (CBOR, JSON, CSV), or as WAV, JPG, PNG, AVI or MP4 files. We also support uploading image datasets with labels in various formats. When you include labels during upload, we attempt to convert your dataset into a format recognized by Studio. here.

Upload mode

- ○ Select individual files
- ◉ Select a folder

Select files

Choose Files No file chosen

Upload into category

- ◉ Automatically split between training and testing
- ○ Training
- ○ Testing

Label

- ○ Infer from filename
- ○ Leave data unlabeled
- ◉ Enter label:

silent

Upload output

```
[2383/2400] Uploading 133400silent.wav OK
[2384/2400] Uploading 122600silent.wav OK
[2385/2400] Uploading 199400silent.wav OK
[2386/2400] Uploading 239100silent.wav OK
[2387/2400] Uploading 25400silent.wav OK
[2388/2400] Uploading 204700silent.wav OK
[2389/2400] Uploading 166400silent.wav OK
[2390/2400] Uploading 80900silent.wav OK
[2391/2400] Uploading 66600silent.wav OK
[2392/2400] Uploading 212500silent.wav OK
[2393/2400] Uploading 70400silent.wav OK
[2394/2400] Uploading 99600silent.wav OK
[2395/2400] Uploading 169900silent.wav OK
[2396/2400] Uploading 118000silent.wav OK
[2397/2400] Uploading 196900silent.wav OK
[2398/2400] Uploading 33600silent.wav OK
[2399/2400] Uploading 125600silent.wav OK
[2400/2400] Uploading 170600silent.wav OK

Done. Files uploaded successful: 2378. Files that failed to upload: 22.

Job completed
```

‹ Back

Upload data

Q1.2

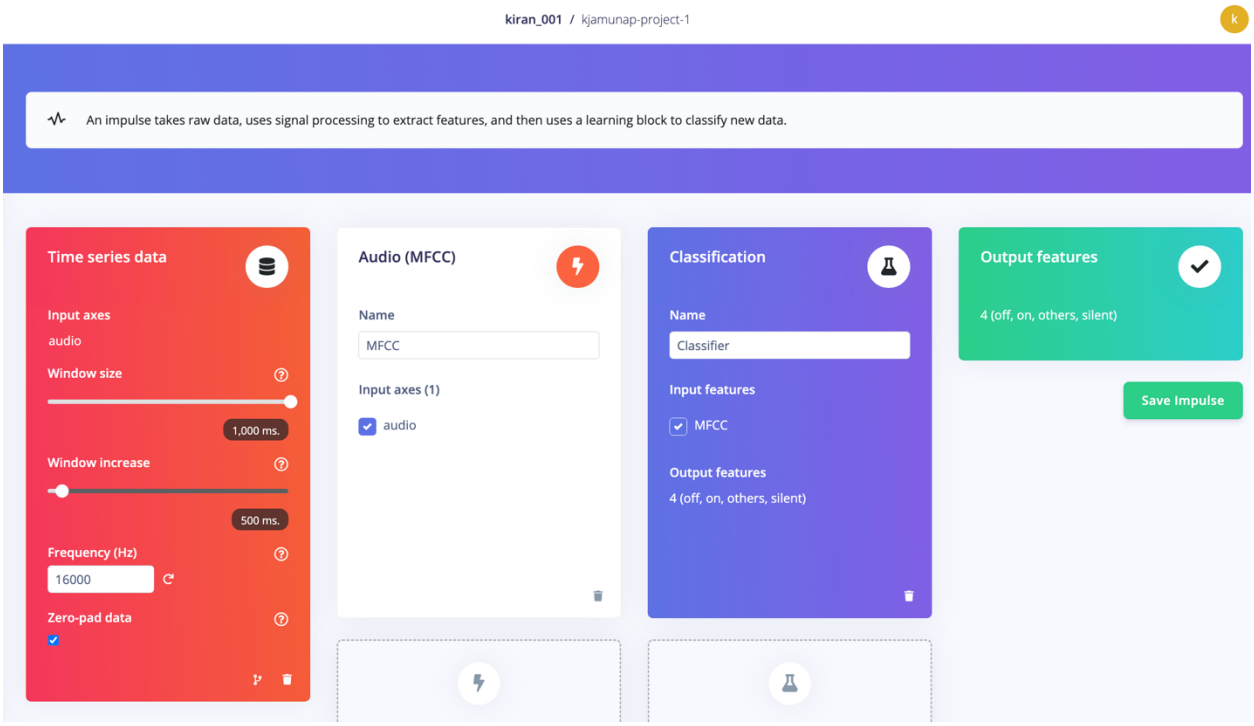
kiran_001 / kjamunap-project-1
An impulse takes raw data, uses signal processing to extract features, and then uses a learning block to classify new data.
Time series data
Input axes
audio
Window size
1,000 ms.
Window increase
500 ms.
Frequency (Hz)
16000
Zero-pad data
Audio (MFCC)
Name
MFCC
Input axes (1)
audio
Classification
Name
Classifier
Input features
MFCC
Output features
4 (off, on, others, silent)
Output features
4 (off, on, others, silent)
Save Impulse

Q1.3

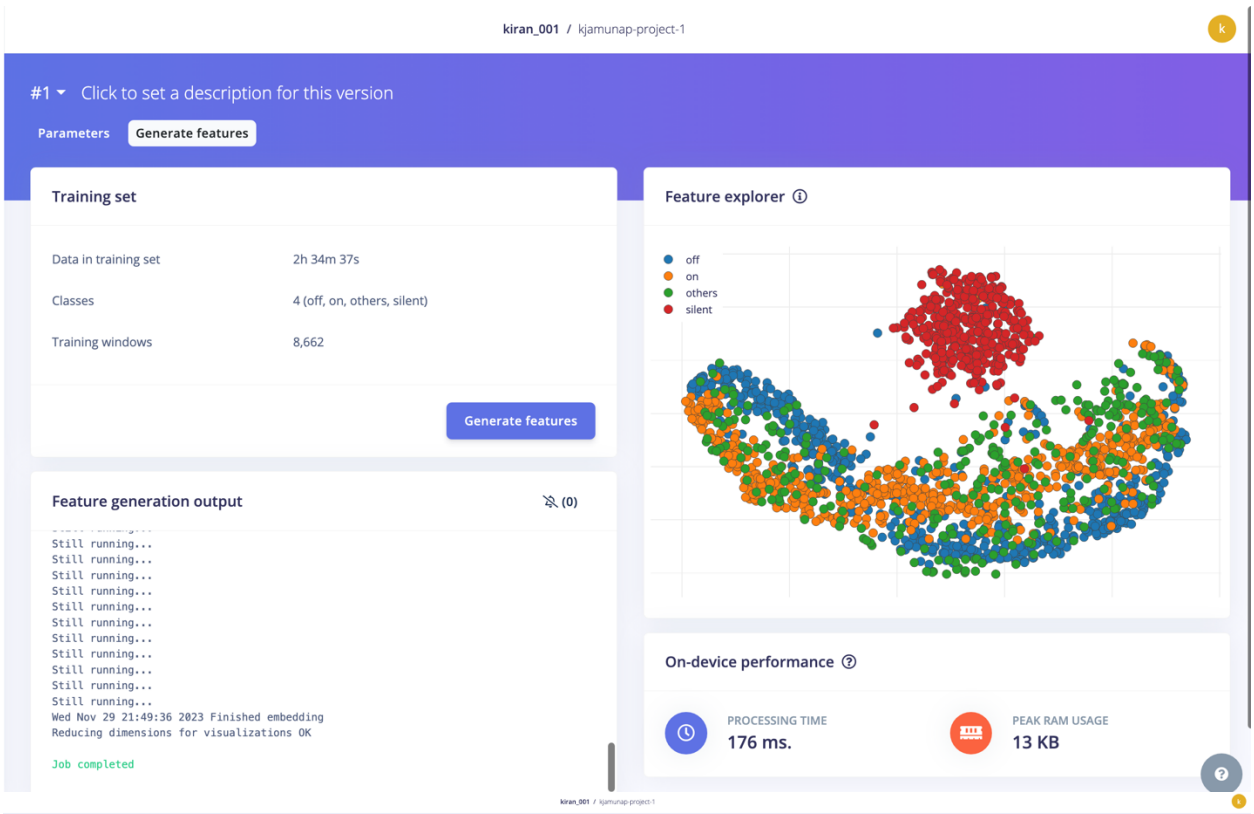
kiran_001 / kjamunap-project-1
#1 Click to set a description for this version
Parameters
Generate features
Training set
Data in training set
2h 34m 37s
Classes
4 (off, on, others, silent)
Training windows
8,662
Generate features
Feature explorer
off
on
others
silent
Feature generation output
(0)
Still running...
Still running...
Still running...
Still running...
Still running...
Still running...
Still running...
Still running...
Still running...
Still running...
Still running...
Wed Nov 29 21:49:36 2023 Finished embedding
Reducing dimensions for visualizations OK
Job completed
On-device performance
PROCESSING TIME
176 ms.
PEAK RAM USAGE
13 KB

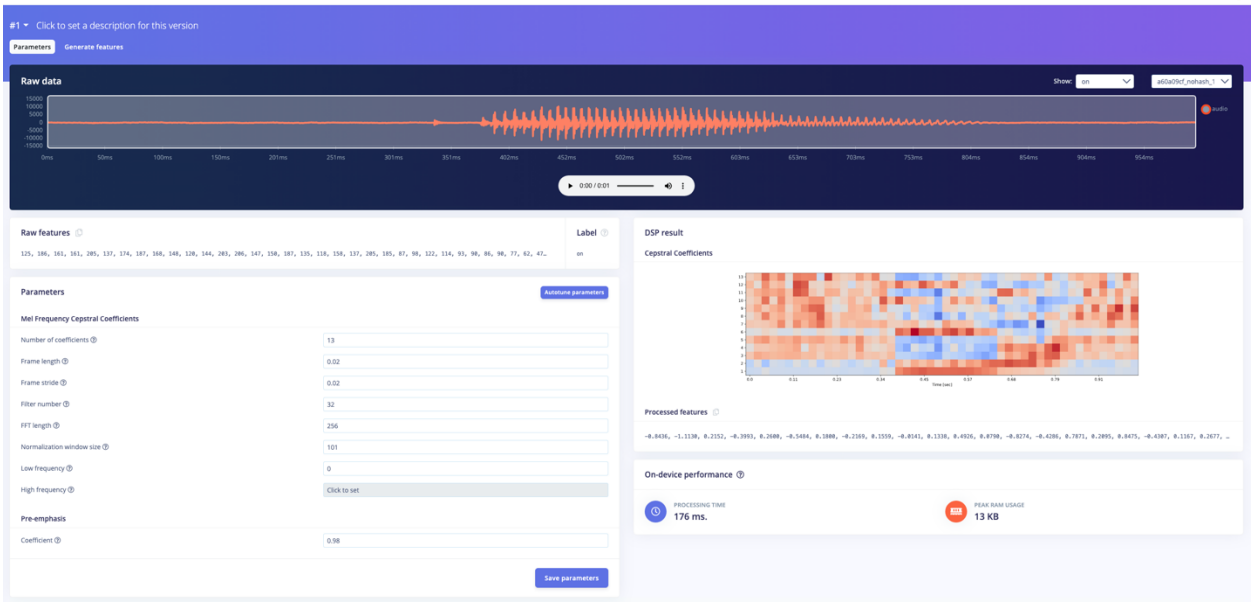
#1 Click to set a description for this version
Parameters
Generate features
Raw data
Raw features
Label
DSP result
Cepstral Coefficients
Parameters
Mel Frequency Cepstral Coefficients
Pre-emphasis
Save parameters
Processed features
On-device performance
PROCESSING TIME
176 ms.
PEAK RAM USAGE
13 KB

Q1.4

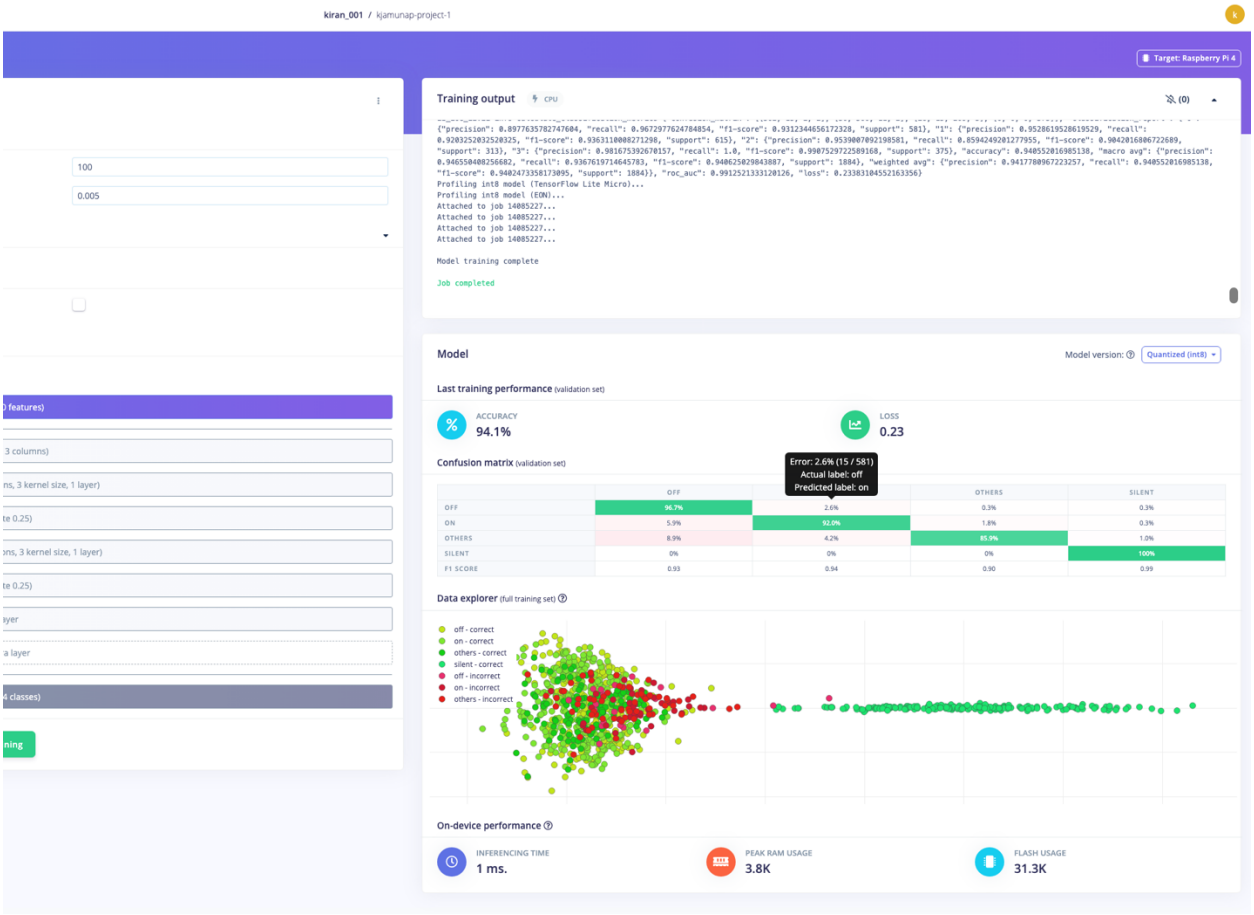
kiran_001 / kjamunap-project-1
Target: Raspberry Pi 4
Training output
CPU
Profiling int8 model (TensorFlow Lite Micro)...
Profiling int8 model (EON)...
Attached to job 14085227...
Attached to job 14085227...
Attached to job 14085227...
Attached to job 14085227...
Model training complete
Job completed
Model
Model version:
Quantized (int8)
Last training performance (validation set)
ACCURACY
94.1%
LOSS
0.23
Confusion matrix (validation set)
Error: 2.6% (15 / 581)
Actual label: off
Predicted label: on
OFF
ON
OTHERS
SILENT
F1 SCORE
Data explorer (full training set)
off - correct
on - correct
others - correct
silent - correct
off - incorrect
on - incorrect
others - incorrect
On-device performance
INFERENCING TIME
1 ms.
PEAK RAM USAGE
3.8K
FLASH USAGE
31.3K

Q2

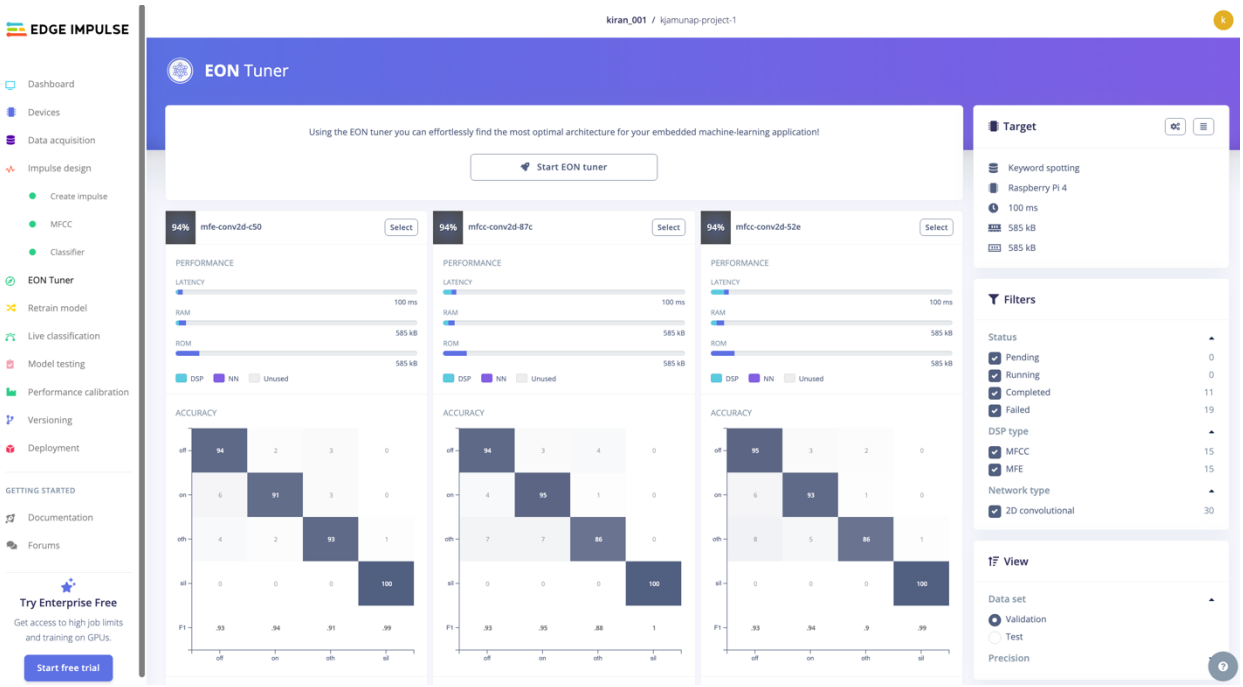
EDGE IMPULSE
Dashboard
Devices
Data acquisition
Impulse design
Create impulse
MFCC
Classifier
EON Tuner
Retrain model
Live classification
Model testing
Performance calibration
Versioning
Deployment
GETTING STARTED
Documentation
Forums
Try Enterprise Free
Get access to high job limits and training on GPUs.
Start free trial
kiran_001 / kjamunap-project-1
EON Tuner
Using the EON tuner you can effortlessly find the most optimal architecture for your embedded machine-learning application!
Start EON tuner
94% mfe-conv2d-c50 Select
94% mfcc-conv2d-87c Select
94% mfcc-conv2d-52e Select
PERFORMANCE
LATENCY
RAM
ROM
100 ms
585 kB
DSP NN Unused
ACCURACY
off on oth sil
F1
Target
Keyword spotting
Raspberry Pi 4
100 ms
585 kB
585 kB
Filters
Status
Pending 0
Running 0
Completed 11
Failed 19
DSP type
MFCC 15
MFE 15
Network type
2D convolutional 30
View
Data set
Validation
Test
Precision

Best model

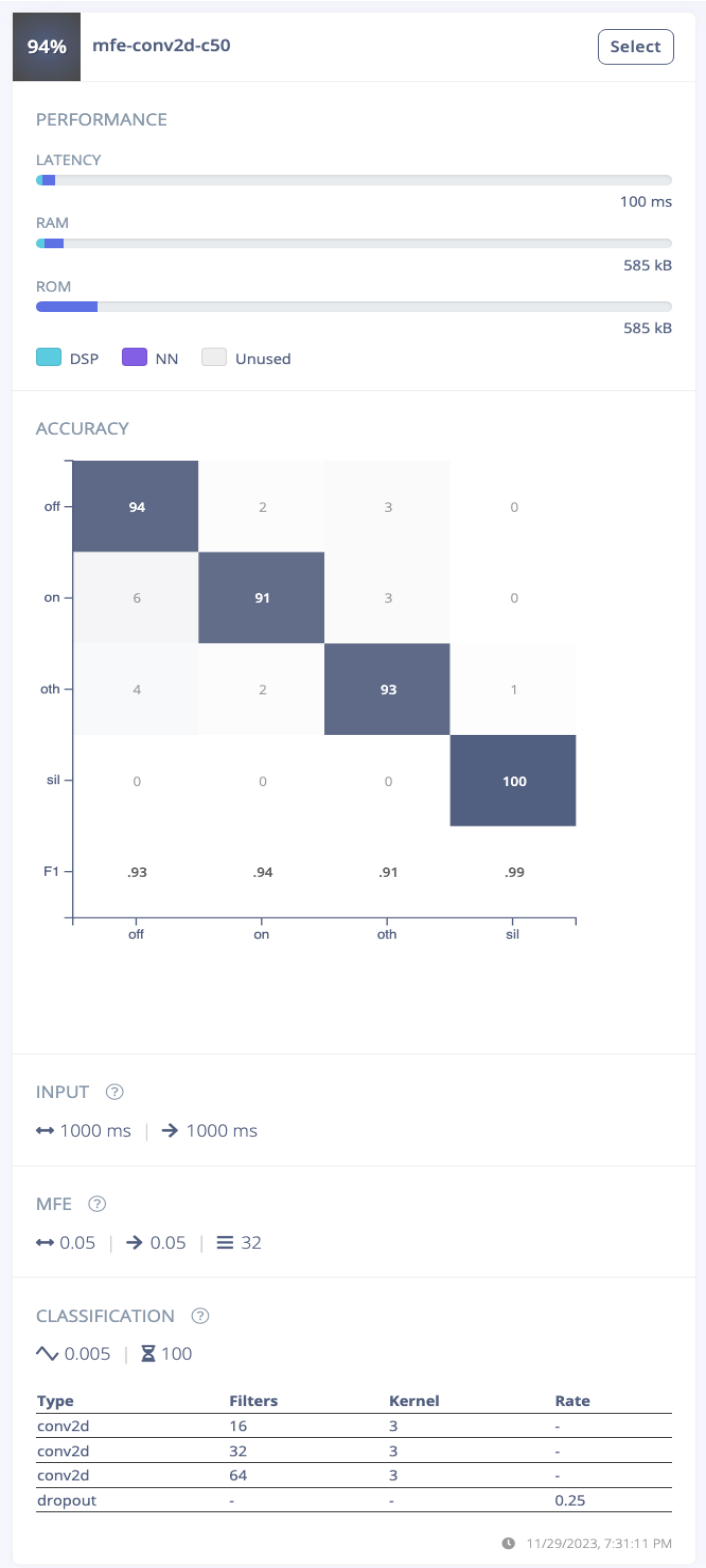
94%
mfe-conv2d-c50
Select
PERFORMANCE
LATENCY
100 ms
RAM
585 kB
ROM
585 kB
DSP
NN
Unused
ACCURACY
off
on
oth
sil
F1
94
2
3
0
6
91
3
0
4
2
93
1
0
0
0
100
.93
.94
.91
.99
off
on
oth
sil
INPUT
1000 ms | 1000 ms
MFE
0.05 | 0.05 | 32
CLASSIFICATION
0.005 | 100
Type
Filters
Kernel
Rate
conv2d
16
3
-
conv2d
32
3
-
conv2d
64
3
-
dropout
-
-
0.25
11/29/2023, 7:31:11 PM

The model is optimized for Raspberry Pi 4 as my latency did not exceed.
My accuracy is around 94%.
Comparing other 2 model which had same accuracy but had bit more latency and the F1 score were less compared to the best model.
So the best model has 94% accuracy, minimum latency, F1 score above 90% for all classes (off - .93, on - .94, others - .91, silent - .91).